

Fig. 5 Latency Distribution per Operation (ML-KEM & ML-DSA)
1 000 trials each | white line = median | SLH-DSA excluded (extreme scale)

